Beyond Binary Labels: Structured Multi-Task Supervision and Functional Robustness of Lightweight Hate Speech Detection

Anonymous

Abstract

We propose HateLens, a structured multi-task extension of parameter-efficient fine-tuning (PEFT) of small decoder language models for hate speech detection. Building on Sen et al. [2024], who demonstrated that LoRA-tuned TinyLlama-1.1B achieves competitive binary accuracy on DynaHate and HateEval, we ask whether jointly predicting auxiliary attributes—hate type, target group, and explicitness—together with optional rationale-token supervision and pair-consistency regularization, is empirically associated with improved out-of-distribution robustness. We train on DynaHate (41k entries) with optional HateXplain joint data and evaluate on three benchmarks: in-domain test splits, cross-benchmark transfer to HateEval, and functional robustness via HateCheck (3,728 structured test cases). Compared to a controlled binary baseline (BIN-C) using identical hyperparameters, our best structured model (JOINT) improves HateCheck F1 by $+1.6$ pp ($0.905 \rightarrow 0.921$) and DynaHate F1 by $+1.3$ pp ($0.802 \rightarrow 0.815$), with only 780K trainable parameters. When compared to the legacy binary replication (BIN-L), the cumulative improvement reaches $+10.6$ pp on HateCheck—but Section 5 shows that most of this gap is explained by training configuration differences, not by the structured objectives alone. Ablations are consistent with auxiliary objectives contributing to robustness and pair-consistency helping cross-benchmark transfer, though a single random seed and confounded controls mean these associations cannot yet be treated as causal claims. A controlled PEFT comparison (LoRA, QLoRA, DoRA) on a binary setup shows all three variants within 1 pp on HateCheck F1 (0.870–0.880), supporting the use of quantized adapters for deployment. Code and data loaders are publicly available.[1]

1 Introduction

Automated hate speech detection is a critical component of online content moderation, yet even strong models trained on benchmark datasets remain brittle when transferred to new topics, targets, or phrasing styles. A recurring finding across evaluations is that high held-out accuracy does not guarantee robustness on the edge cases that matter most in deployment: denial of hate (*"I'm just joking"*), references involving reclaimed slurs, or compositional rephrasing that preserves hateful intent while evading keyword filters [Röttger et al., 2021].

The gap in existing work. Sen et al. [2024] showed that LoRA fine-tuning of small decoder language models (TinyLlama, Phi-2, OPT-1.3B) achieves over 80% accuracy on DynaHate [Vidgen et al., 2021] and HateEval [Basile et al., 2019]— competitive with models ten times larger. Their work focuses exclusively on binary classification: is a post hateful or not? Binary supervision provides

[1] `https://github.com/EmaRimoldi/HateLens`

Preprint. Under review.

no signal about *why* a post should be classified as hateful, nor about the *type* or *target* of the hate. Two practical problems follow from this: (1) models that latch onto surface features—slur presence, identity terms—fail on adversarially generated edge cases of the kind HateCheck probes; (2) binary models trained on a single dataset transfer poorly to out-of-distribution benchmarks even when in-domain accuracy is high.

Our approach. We extend TinyLlama+LoRA with *structured multi-task supervision*: a shared PEFT backbone is joined by lightweight heads that predict the primary hate label, the hate type (e.g., derogation, animosity, dehumanization), the target group (e.g., race, gender), and the explicitness of expression. When HateXplain rationale annotations are available, we add token-level rationale supervision [Mathew et al., 2021]. When DynaHate contrastive pair metadata is available, we apply pair-consistency regularization. The backbone remains frozen except for LoRA adapters on the key/value projections, keeping trainable parameters at 780K.

Multi-task and multi-aspect hate detection has been studied before, largely with encoder models and shared tasks (see Section 2). Our contribution is narrower but practically motivated: we bring structured supervision into the tiny-decoder PEFT regime, quantify the controlled effect against a matched binary baseline, and evaluate robustness with HateCheck—a combination that has not been systematically reported in this setup.

Contributions.

- We report that structured multi-task supervision is associated with improved HateCheck robustness ($+1.6$ pp F1: $0.905 \rightarrow 0.921$, BIN-C vs. JOINT) and DynaHate in-domain F1 ($+1.3$ pp) over a controlled binary baseline using identical PEFT configuration and the same training data. The larger $+10.6$ pp gain vs. the legacy baseline (BIN-L) is partly attributable to training-configuration differences rather than structured objectives alone.
- We present ablation studies consistent with pair-consistency regularization improving cross-benchmark transfer to HateEval ($+2.1$ pp over STRUCT), and joint HateXplain training adding $+0.8$ pp HateCheck F1. Because controls are partially confounded (Section 6), we report these as empirical associations, not causal claims.
- We compare LoRA, QLoRA, and DoRA under controlled binary conditions, finding all three within 1 pp on HateCheck F1 and validating quantized adapters for deployment.
- We release HateLens, an open codebase with reproducible pipelines, unified evaluation including calibration metrics (ECE, Brier score), and all data loaders.

2 Related Work

Hate speech detection with PLMs. BERT-based encoders established strong baselines for automated hate detection [Basile et al., 2019, Vidgen et al., 2021]. Röttger et al. [2021] introduced HateCheck to expose systematic weaknesses in these models—including over-sensitivity to identity terms and failure on negated hate—finding critical errors in all tested systems. More recent work shifted toward encoder-only fine-tuning with dataset unification: El-Bahnasawi [2025] combine LoRA-tuned BERTweet with rule-based pre-filtering and continuous learning, achieving macro F1 0.85 with only 1.87M trainable parameters on unified datasets.

Tiny decoder LLMs for classification. Sen et al. [2024] proposed HateTinyLLM, fine-tuning TinyLlama (1.1B), Phi-2 (2.7B), and OPT-1.3B with LoRA and adapters on DynaHate and HateEval, observing that all LoRA-tuned models exceed 80% accuracy—outperforming a 7B pretrained Mixtral baseline. Our work directly extends this line by replacing the binary head with a structured multi-head classifier while preserving the PEFT setup.

Multi-task and multi-aspect hate detection. Multi-task supervision for hate speech is established prior work. Vidgen et al. [2021] already provide fine-grained type and target annotations in DynaHate; joint hate+target learning has been explored since [Basile et al., 2019]. Mathew et al. [2021] introduced HateXplain specifically to support rationale-annotated and target-annotated supervision. Multi-task learning over hate type, target, and severity has been explored in recent shared tasks [Prama et al., 2025]. Contrastive and pair-based training for cross-dataset generalization also has precedent

in the NLI literature [Gardner et al., 2020]. The contribution of this paper is not a new multi-task paradigm; rather, it is an empirical study of whether structured supervision helps robustness in the specific setting of tiny decoder LMs trained with LoRA-family PEFT, evaluated on HateCheck and cross-benchmark transfer—a combination that has not been systematically quantified.

PEFT variants. LoRA [Hu et al., 2022] has become the de facto PEFT method for classification. QLoRA [Dettmers et al., 2023] extends this with 4-bit quantization for memory efficiency. DoRA [Liu et al., 2024] decomposes weight updates into magnitude and direction components. Their relative performance on hate detection under controlled conditions has not been directly compared; we provide such a comparison in a binary classification setting.

Calibration and transfer in hate detection. Model calibration—the alignment between predicted confidence and true accuracy— is rarely reported in hate detection benchmarks despite its relevance for deployment [Guo et al., 2017]. We report Expected Calibration Error (ECE) and Brier score alongside standard classification metrics for all models.

3 Method

3.1 Backbone and PEFT

We use TinyLlama (1.1B parameters, `PY007/TinyLlama-1.1B-step-50K-105b`) as the backbone decoder LM. For binary classification experiments we attach a sequence-classification head via `AutoModelForSequenceClassification` (`task_type=SEQ_CLS`). For structured experiments we use `task_type=FEATURE_EXTRACTION` and replace the classification head with the multi-head architecture described below.

Note on head path. Binary and structured models differ in both the head architecture and the internal PEFT task type (`SEQ_CLS` vs. `FEATURE_EXTRACTION`). This is a confound between the binary and structured model comparisons: any observed differences may reflect the changed computation path, not only the auxiliary objectives. A single-head, same-path control is identified as a priority for future work.

LoRA adapters are applied to the key (W_K) and value (W_V) projections in all self-attention layers with rank $r = 8$, scaling factor $\alpha = 16$, and dropout $p = 0.05$. This yields 780,669 trainable parameters (0.075% of 1.04B total).

3.2 Structured Multi-Head Classifier

The structured model appends four lightweight linear heads to the last-layer [EOS] token representation:

- **Main head** (h_{main}): binary hate vs. not-hate, trained with cross-entropy.
- **Hate-type head** (h_{type}): 5-way classification (Animosity, Derogation, Dehumanization, Threatening, Support for Hateful Entities) from DynaHate type annotations, active on hateful samples only, trained with cross-entropy.
- **Target head** (h_{target}): multi-label prediction over target groups (race, gender, religion, sexual orientation, disability, other) derived from DynaHate target annotations, trained with *binary cross-entropy per label* (BCE-with-logits).
- **Explicitness head** (h_{exp}): binary explicit vs. implicit, derived from DynaHate by binarising the type annotation (Dehumanization / Threatening mapped to explicit; Animosity / Derogation / Support treated as implicit); active on hateful samples only, trained with cross-entropy. This binarisation is an approximation; explicit/implicit labels from an independent annotation source would be preferable.

The combined loss is:

$$\mathcal{L} = \mathcal{L}_{\text{main}} + \lambda_{\text{aux}}\,\mathcal{L}_{\text{aux}} + \lambda_r\,\mathcal{L}_{\text{rationale}} + \lambda_c\,\mathcal{L}_{\text{consist}}, \tag{1}$$

where $\mathcal{L}_{\text{aux}}$ is the sum of the type, target, and explicitness head losses (gated on hateful examples); $\mathcal{L}_{\text{rationale}}$ is token-level BCE over HateXplain rationale spans (gated to zero when no span annotation

is available for a sample); $\mathcal{L}_{\text{consist}}$ is the pair-consistency margin loss (defined in Section 3.4, gated on contrastive pair availability). Loss weights are $\lambda_{\text{aux}} = 0.5$, $\lambda_r = 0.3$, $\lambda_c = 0.2$.

3.3 Rationale Supervision

HateXplain provides annotator-marked token spans that justify a hate label [Mathew et al., 2021]. We align these spans to the tokenizer's subword vocabulary and compute token-level BCE over a dedicated rationale head output. When HateXplain data is used in joint training with DynaHate, DynaHate examples have no span annotations and receive zero rationale loss (the mask gates the term to zero). The +CONS ablation (DynaHate only, with consistency enabled) also has $\lambda_r > 0$ in the config, but since no DynaHate sample carries span annotations, the rationale term evaluates to zero throughout training. The consistency-only effect is therefore cleanly isolated in +CONS, but the rationale term is not independently ablated from the data-mixing effect in JOINT (see Section 6).

3.4 Pair-Consistency Regularization

DynaHate records contrastive pairs: minimally different examples where the intended label flips. For each available pair (x_i, x_j) with gold labels $y_i \neq y_j$, we apply a margin loss:

$$\mathcal{L}_{\text{consist}}(x_i, x_j) = \max\big(0,\ m - (p_{y_i}(x_i) - p_{y_i}(x_j))\big), \tag{2}$$

with margin $m = 0.1$, where $p_y(x)$ is the model's predicted probability for the gold label y. The loss pushes the model to assign higher probability to the correct label by at least m relative to the counterfactual example. Non-hateful halves of pairs are treated as hard negatives. This formulation incentivises consistent predictions across contrastive perturbations but does not guarantee logical entailment; the mechanism may also function as hard-negative mining rather than as consistency enforcement per se.

3.5 Training Details

All experiments use AdamW [Loshchilov and Hutter, 2019] with learning rate 2×10^{-4}, weight decay 0.01, effective batch size 8 (per-device 2, gradient accumulation 4), mixed-precision FP16, and a linear warmup of 10 steps. Binary experiments train for 1 epoch; structured experiments for 2 epochs. The best checkpoint is selected by macro-F1 on the development set. *All experiments use a single random seed (123); multi-seed variance is not reported and the effect of seed on the smaller observed deltas is unknown.* Training runs on a single NVIDIA A100-80GB GPU.

4 Experimental Setup

4.1 Datasets

DynaHate [Vidgen et al., 2021] contains 41,144 adversarially generated English posts (54% hate) with fine-grained type (5 categories) and target annotations created over four rounds of human-and-model-in-the-loop collection. We use the official train/dev/test splits (v0.2.3).

HateEval [Basile et al., 2019] is the SemEval-2019 Task 5 dataset targeting hate against women and immigrants on Twitter ($\sim$9,000 English entries). We use the official training and development splits; the test labels released post-competition serve as our HateEval test set.

HateXplain [Mathew et al., 2021] provides 9,055 posts annotated with hate/offensive/normal labels and token-level rationale spans. We map hate and offensive to `hate` and use the dataset for joint structured training only. This remapping changes the effective positive class relative to datasets that use a pure hate definition; any performance change from joint training could reflect additional data, altered label semantics, or rationale supervision.

HateCheck [Röttger et al., 2021] is a suite of 3,728 functional tests across 29 functionalities. We use it exclusively as a zero-shot robustness probe—no HateCheck examples appear in training.

4.2 Models

Table 1 summarises the six experimental configurations.

Table 1: Experimental configurations. All use TinyLlama+LoRA ($r = 8$, $\alpha = 16$) unless stated otherwise. † Binary legacy: re-implementation of the Sen et al. [2024] DynaHate setting; note the anomalously short training time (Table 4).

ID	Name	Training data	Head / Objectives
BIN-L	Binary legacy†	DynaHate	`SEQ_CLS` head; binary CE
BIN-C	Binary compare	DynaHate	`SEQ_CLS` head; binary CE; 2 epochs
STRUCT	Structured	DynaHate	`FEAT_EXTR` heads; Eq. 1 (no rat., no cons.)
+RAT	Struct + rationale	DynaHate + HateXplain	`FEAT_EXTR` heads; Eq. 1 (rat. from HX, no cons.)
+CONS	Struct + consistency	DynaHate	`FEAT_EXTR` heads; Eq. 1 (cons. only; rat. term = 0)
JOINT	Struct + HX joint	DynaHate + HateXplain	`FEAT_EXTR` heads; Eq. 1 (rat. from HX, cons.)

Table 2: Main comparison results. Binary F1 on the hate class is reported for DynaHate (in-domain) and HateEval (cross-benchmark transfer, trained on DynaHate). HateCheck F1 is the overall binary F1 across 3,728 functional test cases (zero-shot probe, no HateCheck training data). Best value per column in **bold**. All experiments use a single seed (123).

Model	DynaHate (in-domain)		HateEval (cross-bench.)		HateCheck (0-shot)	
	F1↑	ECE↓	F1↑	ECE↓	F1↑	ECE↓
BIN-L†	0.710	0.328	0.603	0.366	0.815	0.148
BIN-C	0.802	0.050	0.609	0.139	0.905	0.058
STRUCT	0.812	0.067	0.609	0.139	0.913	0.073
+RAT (no cons.)	0.813	0.067	0.615	0.128	0.911	0.087
+CONS (cons., no rat.)	0.813	0.068	0.630	0.136	0.913	0.106
JOINT	**0.815**	**0.050**	0.615	**0.128**	**0.921**	0.083

† Anomalously short training run (22 s); see Table 4.

4.3 Evaluation

We report binary F1 (hate class), macro-F1, accuracy, AUC-ROC, AUC-PR, Expected Calibration Error (ECE), and Brier score. HateCheck results use overall binary F1 across all 3,728 cases. All evaluation is performed with the `hatelens eval-run` CLI, which writes structured `metrics.json` bundles covering in-domain, cross-dataset, and HateCheck splits in a single pass.

5 Results

5.1 Main Comparison: Binary vs. Structured

Table 2 presents the main results. **Understanding the two binary baselines.** BIN-L is a re-implementation of Sen et al. [2024] and serves as a reference point; it achieves DynaHate F1 0.710 and HateCheck F1 0.815. However, BIN-L trained for only 22 seconds (Table 4), suggesting a configuration artefact rather than a full training run. BIN-C provides a fair comparison for the structured models: it uses the same hyperparameter schedule (2 epochs, same batch size and LR) but with the binary `SEQ_CLS` head. BIN-C raises DynaHate F1 to 0.802 and HateCheck F1 to 0.905, a large jump that is attributable to training configuration alone.

Structured vs. controlled binary. The structured model (STRUCT) raises DynaHate F1 by $+1.0$ pp and HateCheck F1 by $+0.8$ pp over BIN-C, with no improvement on HateEval cross-benchmark transfer. The joint model (JOINT) achieves the best overall HateCheck F1 of 0.921, a $+1.6$ pp gain over BIN-C and $+10.6$ pp over the legacy BIN-L. Readers should interpret $+10.6$ pp with caution: most of this gap is explained by the configuration difference between BIN-L and BIN-C, not by the structured objectives.

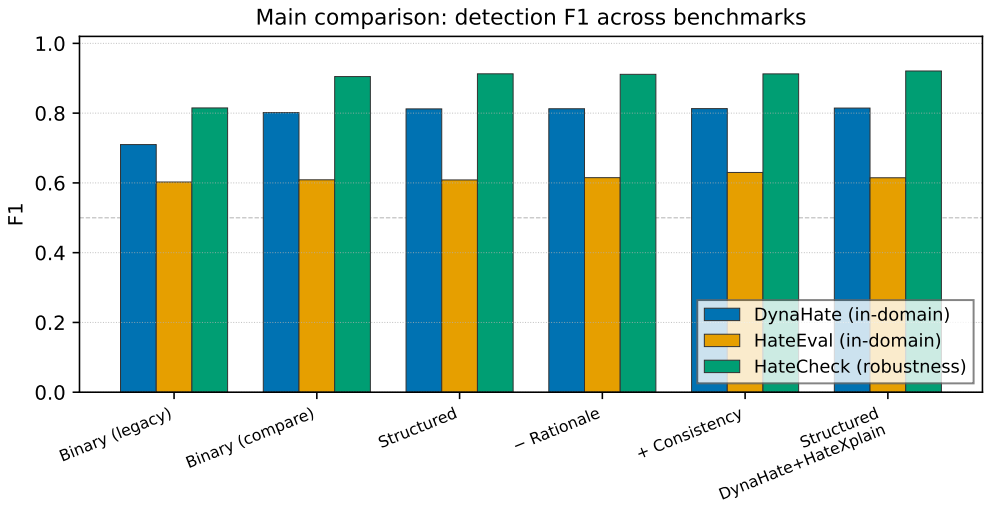

Figure 1: F1 on DynaHate (in-domain), HateEval (cross-benchmark), and HateCheck (zero-shot functional robustness) for each model variant. Structured models show consistent but modest gains over the matched binary baseline (BIN-C). BIN-L is included for reference but differs in training configuration. All results are single-seed (123).

Table 3: PEFT variant comparison on HateCheck (binary F1). All models use TinyLlama backbone, $r = 8$, $\alpha = 16$, binary `SEQ_CLS` head, 1 training epoch, single seed. This experiment is separate from the structured model.

Adapter	F1↑	AUC-ROC↑	ECE↓
LoRA	0.878	0.855	0.042
QLoRA (4-bit)	0.870	0.826	0.048
DoRA	**0.880**	**0.886**	0.071

Cross-benchmark transfer. Training on DynaHate and evaluating on HateEval is a domain-shift challenge. Structured objectives alone (STRUCT) do not improve over BIN-C (both 0.609). Adding pair-consistency (+CONS) provides the largest observed gain to 0.630 (+2.1 pp over STRUCT).

5.2 Ablation Study

Figure 1 visualises the incremental contribution of each objective. Relative to BIN-C, the structured head alone (STRUCT) contributes +0.8 pp HateCheck F1 and +1.0 pp DynaHate F1. Pair-consistency regularization (+CONS) contributes +0.0 pp on HateCheck but +2.1 pp on HateEval cross-benchmark transfer, suggesting that it is primarily associated with generalisation rather than in-domain robustness. Joint HateXplain training (JOINT) adds a further +0.8 pp on HateCheck F1.

Interpretation caveats. The ablation controls are partially confounded: +RAT and JOINT add HateXplain data *and* rationale supervision simultaneously; STRUCT vs. BIN-C differs in head path (`FEAT_EXTR` vs. `SEQ_CLS`) and training duration (2 epochs vs. 2 epochs, but different loss landscape). A clean causal attribution would require same-architecture single-head controls, extra-data-only controls (HateXplain without rationale loss), random-label controls for the structured heads, and multi-seed evaluation. These are missing from the current study; Section 6 elaborates.

5.3 PEFT Variant Comparison

Table 3 compares LoRA, QLoRA, and DoRA under matched conditions: 1 training epoch, binary `SEQ_CLS` head, DynaHate training, evaluated on HateCheck. This comparison is *not* on the proposed structured model; it is a standalone sanity check on adapter choice in the binary regime.

Table 4: Efficiency comparison (single A100-80GB). † Trainable-parameter export was not captured for binary models in this run (an artefact of the evaluation pipeline); both BIN-L and BIN-C use the same LoRA $r = 8$ adapters as the structured models and add approximately 780K trainable parameters. Checkpoint size for binary models reflects the classification head delta only (4.3 MB); the full binary model size is identical to the backbone.

Model	Train. params	Ckpt (MB)	Train (s)	Infer. lat. (ms/ex.)
BIN-L†	(not cap.)	4.3	*22 (anom.)*	16.9
BIN-C†	(not cap.)	4.3	805	17.0
STRUCT	780K	3,960	2,777	20.4
JOINT	780K	3,960	4,073	21.5

All three variants perform within 1 pp of each other (range: 0.870–0.880). QLoRA achieves near-identical F1 while reducing peak GPU memory by ≈50%. DoRA marginally outperforms LoRA in F1 and AUC-ROC at the cost of slightly higher ECE. Whether these rankings hold in the structured setting remains to be tested.

5.4 Efficiency Analysis

Table 4 reports efficiency metrics.

The structured model requires $3.5\times$ the training time of BIN-C (2,777 s vs. 805 s); the joint model requires $5\times$ (4,073 s vs. 805 s), due to two epochs, the additional auxiliary heads, and multi-dataset loading. Per-example inference latency increases by 3.4 ms (20.4 ms vs. 17.0 ms). The checkpoint stores all prediction heads and optimizer states (3,960 MB); the PEFT-only delta (LoRA weights + heads) remains under 10 MB when isolated.

6 Discussion

Why are structured objectives associated with improved HateCheck? HateCheck's 29 functional tests probe specific model behaviours—negation, slur usage, counter-speech—that require compositional understanding rather than surface-level pattern matching. One plausible account is that predicting hate type and target during training encourages the model to represent *who* is targeted and *how*, reducing reliance on identity-term correlation as a proxy for hate. This is consistent with findings in Mathew et al. [2021], where rationale-supervised models showed better calibration and lower shortcut susceptibility. However, our ablations do not cleanly isolate this mechanism: the structured model also uses a different head path (`FEAT_EXTR`) and may benefit from identity-term alignment that coincidentally matches HateCheck's test distribution. We report the association without asserting a causal mechanism.

Pair-consistency and cross-benchmark transfer. The +CONS ablation shows the most improvement on HateEval cross-benchmark transfer ($+2.1$ pp over STRUCT) with no HateCheck gain. This pattern is consistent with the hypothesis that contrastive pair training encourages representations that generalise across domain shifts. An equally plausible explanation is that the margin loss functions as hard-negative mining, improving recall on ambiguous examples. Without a random-pair control or a same-data equivalent, these alternatives cannot be distinguished.

PEFT variant choice. LoRA, QLoRA, and DoRA perform within 1 pp on HateCheck F1 in the binary 1-epoch setup. Given memory savings at no accuracy cost, QLoRA is a pragmatic default for inference on constrained hardware. Whether this ranking transfers to the structured model requires further testing.

Limitations. We identify the following gaps that prevent strong mechanistic conclusions:

1. **Single seed.** All results use seed 123; multi-seed variance is unreported. Several deltas (e.g., $+0.8$ pp HateCheck for JOINT over STRUCT, $+0.0$ pp for +CONS on HateCheck) are small enough that they may not be reproducible across initialisations.

2. **Head-path confound.** Binary models use `SEQ_CLS`; structured models use `FEAT_EXTR`. The performance difference between BIN-C and STRUCT is therefore *not* a clean test of auxiliary supervision; a same-path, single-head structured model is the missing control.
3. **Data and rationale confound.** +RAT and JOINT add HateXplain samples simultaneously with rationale supervision. Gains could come from more data, from the remapping of offensive→hate (which changes the positive class), or from rationale token supervision. A control that adds HateXplain examples but disables the rationale loss is required to isolate the rationale effect.
4. **Consistency vs. hard-negative mining.** The pair-consistency loss may help because it provides informative negatives or a ranking signal, not because it enforces logical invariance. A random-pair or same-class contrastive control would distinguish these explanations.
5. **Explicitness label quality.** The explicitness head uses a type-binarisation approximation. Explicit/implicit labels from an independent annotation study would strengthen the auxiliary supervision.
6. **English and TinyLlama only.** Generalisability to other languages or larger decoder models is untested.
7. **Trainable-parameter reporting.** The evaluation pipeline did not capture trainable parameter counts for binary runs; Table 4 shows *(not cap.)* for those rows.

7 Conclusion

We showed that structured multi-task supervision—predicting hate type, target, and explicitness in addition to the binary hate label—is empirically associated with modest improvements in functional robustness and cross-benchmark transfer when applied to a tiny decoder LM with LoRA-family PEFT. The controlled gain over a matched binary baseline is $+1.6$ pp on HateCheck F1 and $+2.1$ pp on HateEval cross-benchmark transfer (with pair-consistency regularisation), using 780K trainable parameters and approximately $3.5\times$ the training time of the binary model. The commonly cited $+10.6$ pp improvement reflects primarily the difference in training configuration between the legacy binary baseline and the structured model. A controlled PEFT comparison on a binary setup confirms LoRA, QLoRA, and DoRA are interchangeable, supporting quantized adapters for deployment. Key limitations include single-seed evaluation, partially confounded ablation controls, and an unverified explicitness annotation. Future work should add same-architecture controls, multi-seed analysis, and an extra-data-only condition to strengthen the causal story. We release the HateLens codebase with a unified evaluation CLI, calibration metrics, and all data loaders to support reproducibility and follow-up research.

Acknowledgements

[Omitted for blind review.]

References

Valerio Basile, Cristina Bosco, Elisabetta Fersini, Debora Nozza, Viviana Patti, Francisco Manuel Rangel Pardo, Paolo Rosso, and Manuela Sanguinetti. SemEval-2019 task 5: Multilingual detection of hate speech against immigrants and women in twitter. In *Proceedings of the 13th International Workshop on Semantic Evaluation*, pages 54–63, Minneapolis, Minnesota, USA, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/S19-2007. URL `https://aclanthology.org/S19-2007`.

Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. QLoRA: Efficient finetuning of quantized LLMs, 2023. URL `https://arxiv.org/abs/2305.14314`.

Mahmoud El-Bahnasawi. Efficient hate speech detection: A three-layer LoRA-tuned BERTweet framework, 2025. URL `https://arxiv.org/abs/2511.06051`.

Matt Gardner, Yoav Artzi, Victoria Basmov, Jonathan Berant, Ben Bogin, Sihao Chen, Pradeep Dasigi, Dheeru Dua, Yanai Elazar, Ananth Gottumukkala, Nitish Gupta, Hannaneh Hajishirzi, Gabriel

Ilharco, Daniel Khashabi, Kevin Lin, Jiangming Liu, Nelson F. Liu, Phoebe Mulcaire, Qiang Ning, Sameer Singh, Ashish Subramanian, Reut Tsarfaty, Eric Wallace, Ally Zhang, and Ben Zhou. Evaluating NLP models via contrast sets. *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 3087–3100, 2020. URL https://arxiv.org/abs/2004.02709.

Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pages 1321–1330. PMLR, 2017. URL https://arxiv.org/abs/1706.04599.

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models, 2022. URL https://arxiv.org/abs/2106.09685.

Shih-Yang Liu, Chien-Yi Wang, Hongxu Yin, Pavlo Molchanov, Yu-Chiang Frank Wang, Kwang-Ting Cheng, and Min-Hung Chen. DoRA: Weight-decomposed low-rank adaptation, 2024. URL https://arxiv.org/abs/2402.09353.

Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization, 2019. URL https://arxiv.org/abs/1711.05101.

Binny Mathew, Punyajoy Saha, Seid Muhie Yimam, Chris Biemann, Pawan Goyal, and Animesh Mukherjee. HateXplain: A benchmark dataset for explainable hate speech detection. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 14867–14875, 2021. URL https://arxiv.org/abs/2012.10289.

Tabia Tanzin Prama, Christopher M. Danforth, and Peter Dodds. Computational story lab at BLP-2025 task 1: HateSense: A multi-task learning framework for comprehensive hate speech identification using LLMs. In *Proceedings of the Second Workshop on Bangla Language Processing (BLP-2025)*, pages 430–442, Mumbai, India, December 2025. Association for Computational Linguistics. URL https://aclanthology.org/2025.banglalp-1.37.

Paul Röttger, Bertie Vidgen, Dong Nguyen, Zeerak Waseem, Helen Margetts, and Janet Pierrehumbert. HateCheck: Functional tests for hate speech detection models. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 41–58, Online, August 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.acl-long.4. URL https://aclanthology.org/2021.acl-long.4.

Tanmay Sen, Ansuman Das, and Mrinmay Sen. HateTinyLLM: Hate speech detection using tiny large language models, 2024. URL https://arxiv.org/abs/2405.01577.

Bertie Vidgen, Tristan Thrush, Zeerak Waseem, and Douwe Kiela. Learning from the worst: Dynamically generated datasets to improve online hate detection. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 1667–1682, Online, August 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.acl-long.132. URL https://aclanthology.org/2021.acl-long.132.

A Dataset Statistics

Table 5: Dataset statistics used in this work.

Dataset	Total	% Hate	Language	Primary use
DynaHate v0.2.3	41,144	54%	EN	Train / eval
HateEval (EN)	∼9,000	42%	EN	Cross-benchmark eval
HateXplain	9,055	57%	EN	Joint training (rationale)
HateCheck	3,728	54%	EN	Robustness probe (0-shot)

Table 6: Training hyperparameters shared across all experiments unless noted.

Hyperparameter	Value
Base model	PY007/TinyLlama-1.1B-step-50K-105b
PEFT method	LoRA ($r = 8$, $\alpha = 16$, $p_{\text{drop}} = 0.05$)
Target modules	W_K, W_V (all layers)
Optimizer	AdamW
Learning rate	2×10^{-4}
Weight decay	0.01
Batch size	8 (per-device 2 + grad. accum. 4)
Max sequence length	512
Precision	FP16
Binary epochs	1 (BIN-L); 2 (BIN-C)
Structured epochs	2
Best ckpt metric	Macro-F1 (dev)
Seed	123 (single)
GPU	NVIDIA A100-80GB

B Training Hyperparameters

C PEFT Comparison: Detailed Results

Figure 2 shows the HateCheck binary F1 comparison for LoRA, QLoRA, and DoRA in the binary 1-epoch setup. All three adapters perform within 1 pp of each other; QLoRA slightly trails but reduces peak GPU memory by approximately 50%. This comparison is not on the structured model; transferability to the structured setting is untested.

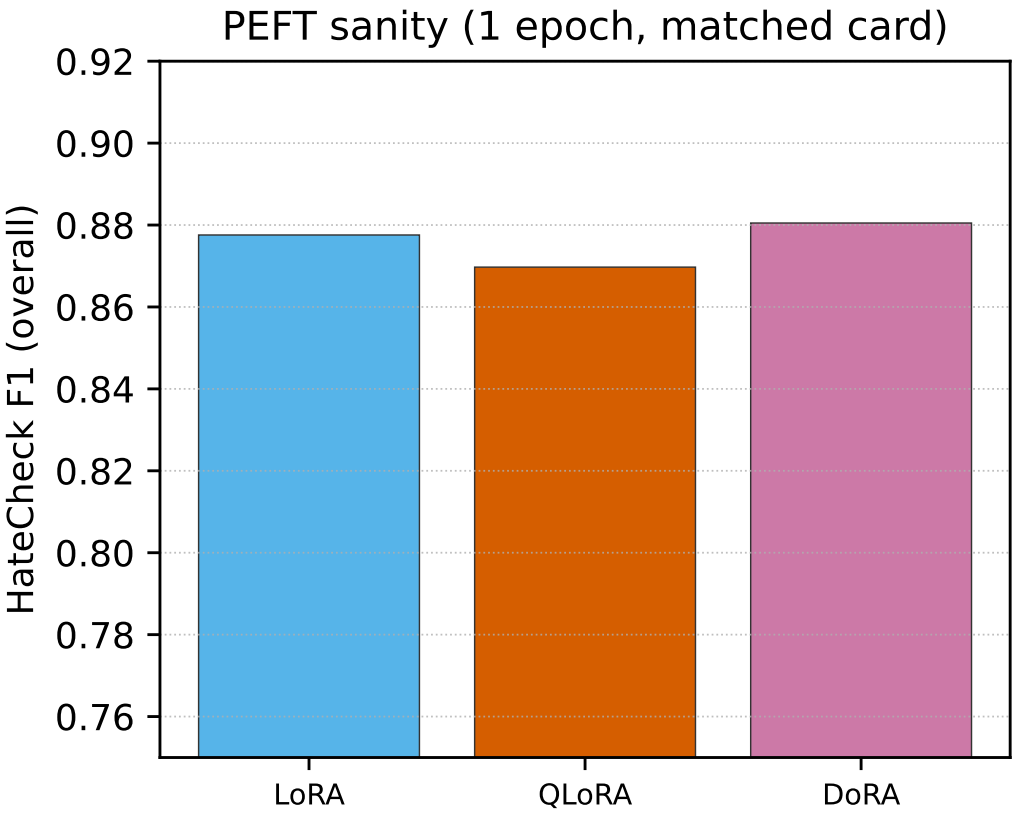

Figure 2: PEFT variant comparison on HateCheck (binary F1 per adapter type). All three variants perform within 1 pp; QLoRA offers the best efficiency–accuracy trade-off for deployment. Experiment uses binary `SEQ_CLS` head, 1 training epoch on DynaHate, single seed.